

EEG Classification Model

Project Overview:

In this project, I built a classification model to analyze EEG data and classify it into different categories. EEG data is widely used in neuroscience and medical fields, including the diagnosis of epilepsy. I used two EEG datasets to train and evaluate my model.

Dataset

The dataset used contains EEG recordings from patients with epilepsy. It includes various seizure types and non-seizure data Link for the Dataset: <https://physionet.org/content/chbmit/1.0.0/>

1. Data Preprocessing:

Function to read a single EDF file, extract, and split data into non-seizure and seizure segments

```
In [86]: import os
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns

from scipy.fft import fft, rfft, rfftfreq
from scipy.stats import skew, kurtosis
from scipy.signal import welch
from sklearn.model_selection import train_test_split
from sklearn.metrics import accuracy_score, classification_report, confusion_matrix, roc_curve, auc
from sklearn.preprocessing import MinMaxScaler

from keras.models import Sequential
from keras.layers import Dense, Dropout, Conv1D, MaxPooling1D, Flatten, Bidirectional, LSTM

import pyedflib
import pywt
import warnings
warnings.filterwarnings('ignore')
```

```

In [2]: def split_edf(file_path, seizure_start, seizure_end, pre_seizure_duration=10):
        with pyedflib.EdfReader(file_path) as edf:
            # Determine the sampling rate
            sampling_rate = edf.getSampleFrequencies()[0]

            # Calculate the indices for the non-seizure and seizure data
            pre_seizure_end_idx = int(seizure_start * sampling_rate)
            pre_seizure_start_idx = pre_seizure_end_idx - int(pre_seizure_duration * sampling_rate)
            seizure_end_idx = int(seizure_end * sampling_rate)

            # Read and split data for each channel
            non_seizure_data = [edf.readSignal(i)[pre_seizure_start_idx:pre_seizure_end_idx] for i in range(edf.signals_in_file)]
            seizure_data = [edf.readSignal(i)[pre_seizure_end_idx:seizure_end_idx] for i in range(edf.signals_in_file)]

            # Extract signal labels
            signal_labels = edf.getSignalLabels()

        return non_seizure_data, seizure_data, signal_labels

file_path = '/Users/archie/Desktop/FDA Project 3/DATASET/chb-mit-scalp-eeG-database-1.0.0/chb01/chb01_03.edf'

seizure_start_time = 2996
seizure_end_time = 3036

# Split the file into non-seizure and seizure parts
non_seizure_data, seizure_data, signal_labels = split_edf(file_path, seizure_start_time, seizure_end_time)

# Convert the split data into pandas DataFrames
df_non_seizure = pd.DataFrame(np.transpose(non_seizure_data), columns=signal_labels)
df_seizure = pd.DataFrame(np.transpose(seizure_data), columns=signal_labels)

```

```

In [3]: # Replace with the paths where you want to save the CSV files
non_seizure_csv_path = '/Users/archie/Desktop/FDA Project 3 Data/Seizure and Non-Seizure/chb01/03non_seizure_data.csv'
seizure_csv_path = '/Users/archie/Desktop/FDA Project 3 Data/Seizure and Non-Seizure/chb01/03seizure_data.csv'

# Save the data to CSV
df_non_seizure.to_csv(non_seizure_csv_path, index=False)
df_seizure.to_csv(seizure_csv_path, index=False)

```

Combining all the Seizure and Non Seizure Csv

```

In [4]: input_path = '/Users/archie/Desktop/FDA Project 3 Data/Seizure and Non-Seizure'

def concat_data(input_path):
    folders = os.listdir(input_path)
    df_seizure = pd.DataFrame()
    df_nonseizure = pd.DataFrame()
    i,j=0,0

    for folder in folders:
        if 'chb' in folder:
            files=os.listdir(input_path+'//'+folder)
            for file in files:
                file=file.lower()
                extension=os.path.splitext(file)[1]
                if extension=='.csv':
                    if 'non_seizure_data' in file:
                        data=pd.read_csv(input_path+'//'+folder+'//'+file)
                        df_nonseizure = pd.concat([df_nonseizure,data],ignore_index=True)
                        i+=1
                    else:
                        data=pd.read_csv(input_path+'//'+folder+'//'+file)
                        df_seizure=pd.concat([df_seizure,data],ignore_index=True)
                        j+=1

    print('The number of seizure files:',i)
    print('The number of non seizure files:',j)
    return df_seizure,df_nonseizure

df_seizure,df_nonseizure=concat_data(input_path)

df_seizure['seizure']=1
df_nonseizure['seizure']=0

df_nonseizure.to_csv(input_path+'//'+non_seizure_data.csv')
df_seizure.to_csv(input_path+'//'+seizure_data.csv')

ch_labels = ['FP1-F7', 'C3-P3', 'C4-P4', 'CZ-PZ', 'F3-C3', 'F4-C4', 'F7-T7',
             'F8-T8', 'FP1-F3', 'FP2-F4', 'FP2-F8', 'FT10-T8', 'FT9-FT10', 'FZ-CZ',
             'P3-O1', 'P4-O2', 'P7-O1', 'P7-T7', 'P8-O2', 'T7-FT9', 'T7-P7',
             'T8-P8-0', 'T8-P8-1', 'seizure']

```

The number of seizure files: 196
The number of non seizure files: 196

```
In [6]: # check columns in ch_labels
        for cnt in ch_labels:
            if cnt not in df_seizure.columns:
                print(cnt)
```

T8-P8-0
T8-P8-1

```
In [7]: # check columns in ch_labels
        for cnt in ch_labels:
            if cnt not in df_nonseizure.columns:
                print(cnt)
```

T8-P8-0
T8-P8-1

```
In [8]: # discarding irrelevant channels from df_seizure
df_seizure=df_seizure[df_seizure.columns[df_seizure.columns.isin(ch_labels)]]
```

```
In [9]: # discarding irrelevant channels from df_seizure
df_nonseizure=df_nonseizure[df_nonseizure.columns[df_nonseizure.columns.isin(ch_labels)]]
```

```
In [10]: df_seizure.isnull().sum()
```

```
Out[10]: FP1-F7      124416
          F7-T7      124416
          T7-P7      124416
          P7-O1      124416
          FP1-F3      124416
          F3-C3      124416
          C3-P3      124416
          P3-O1      124416
          FZ-CZ      124416
          CZ-PZ      124416
          FP2-F4      124416
          F4-C4      124416
          C4-P4      124416
          P4-O2      124416
          FP2-F8      124416
          F8-T8      124416
          P8-O2      124416
          P7-T7      152576
          T7-FT9      152576
          FT9-FT10    152576
          FT10-T8     152576
          seizure      0
          dtype: int64
```

```
In [11]: df_nonseizure.isnull().sum()
```

```
Out[11]: FP1-F7      33280
          F7-T7      33280
          T7-P7      33280
          P7-O1      33280
          FP1-F3      33280
          F3-C3      33280
          C3-P3      33280
          P3-O1      33280
          FZ-CZ      33280
          CZ-PZ      33280
          FP2-F4      33280
          F4-C4      33280
          C4-P4      33280
          P4-O2      33280
          FP2-F8      33280
          F8-T8      33280
          P8-O2      33280
          P7-T7      43520
          T7-FT9      43520
          FT9-FT10     43520
          FT10-T8      43520
          seizure      0
          dtype: int64
```

```
In [12]: df_seizure.interpolate(method='linear', axis=0, inplace=True)
```

```
In [13]: df_seizure.head()
```

```
Out[13]:
```

	FP1-F7	F7-T7	T7-P7	P7-O1	FP1-F3	F3-C3	C3-P3	P3-O1	FZ-CZ	CZ-PZ	...	C4-
0	31.452991	-100.610501	-2.930403	4.102564	9.181929	-23.247863	-42.393162	-9.963370	5.665446	-57.631258	...	-41.2210
1	40.048840	-100.610501	-12.698413	12.698413	0.195360	-24.420024	-38.485958	3.711844	2.539683	-52.551893	...	-40.0488
2	36.141636	-87.716728	-18.949939	26.373626	-0.586081	-27.155067	-32.625153	18.559219	-2.148962	-37.704518	...	-22.0757
3	43.565324	-84.981685	-2.539683	16.214896	-3.711844	-25.982906	-22.075702	25.982906	-10.354090	-31.452991	...	-22.4664
4	57.631258	-96.312576	4.884005	20.122100	-6.446886	-29.890110	-15.433455	40.830281	-14.652015	-23.247863	...	-8.0097

5 rows × 22 columns


```
In [14]: df_nonseizure.interpolate(method='linear', axis=0, inplace=True)
```

```
In [15]: df_nonseizure.head()
```

```
Out[15]:
```

	FP1-F7	F7-T7	T7-P7	P7-O1	FP1-F3	F3-C3	C3-P3	P3-O1	FZ-CZ	CZ-PZ	...	C4-P4	
0	-29.890110	-1.367521	1.367521	60.757021	1.758242	-9.572650	10.354090	26.764347	3.321123	9.572650	...	0.586081	-1
1	-35.750916	-16.214896	2.539683	70.525031	1.367521	-18.559219	8.400488	27.545788	-4.102564	9.572650	...	0.195360	0
2	-33.797314	-21.294261	1.367521	70.134310	4.102564	-23.247863	9.181929	24.420024	-10.744811	12.307692	...	-3.711844	-2
3	-27.936508	-9.572650	-4.884005	58.412698	8.791209	-25.982906	12.307692	18.168498	-16.996337	16.996337	...	0.195360	-4
4	-27.545788	-17.777778	3.711844	54.114774	4.493284	-24.810745	12.307692	18.949939	-20.512821	18.168498	...	-2.148962	-1

5 rows × 22 columns

2. Feature Extraction:

```
In [16]: # Define a function to calculate time-domain features
def extract_time_domain_features(signal):
    features = {
        'mean': np.mean(signal),
        'std': np.std(signal),
        'variance': np.var(signal),
        'max': np.max(signal),
        'min': np.min(signal),
        'skewness': skew(signal),
        'kurtosis': kurtosis(signal)
    }
    return features

# Define a function to calculate frequency-domain features
def extract_frequency_domain_features(signal, fs):
    freqs, psd = welch(signal, fs=fs)
    bands = {'alpha': (8, 12), 'beta': (12, 30)}

    band_power = {}
    for band, (low_freq, high_freq) in bands.items():
        idx_band = np.logical_and(freqs >= low_freq, freqs <= high_freq)
        band_power[band] = np.trapz(psd[idx_band], freqs[idx_band])

    return band_power
```

```
In [17]: # Assume a sampling frequency of 256Hz for the EEG data
fs = 256
```

```
In [18]: # Apply the feature extraction functions
seizure_time_features = df_seizure.apply(extract_time_domain_features, axis=0)
seizure_freq_features = df_seizure.apply(extract_frequency_domain_features, axis=0, fs=fs)
```

```
In [19]: non_seizure_time_features = df_nonseizure.apply(extract_time_domain_features, axis=0)
non_seizure_freq_features = df_nonseizure.apply(extract_frequency_domain_features, axis=0, fs=fs)
```

```
In [20]: # Combine time and frequency domain features
seizure_features = pd.concat([seizure_time_features, seizure_freq_features], axis=1)
non_seizure_features = pd.concat([non_seizure_time_features, non_seizure_freq_features], axis=1)
```

```
In [21]: # Flatten the dictionaries in each row into separate columns
seizure_time_features = seizure_time_features.apply(pd.Series)
seizure_freq_features = seizure_freq_features.apply(pd.Series)
```

```
In [22]: non_seizure_time_features = non_seizure_time_features.apply(pd.Series)
non_seizure_freq_features = non_seizure_freq_features.apply(pd.Series)
```

```
In [23]: # Combine seizure and non-seizure data and add labels
seizure_features['seizure'] = 1
non_seizure_features['seizure'] = 0
```

```
In [24]: final_df = pd.concat([seizure_features, non_seizure_features], ignore_index=True)
```

```
In [33]: final_df
```

```
Out[33]:
```

	FP1-F7	F7-T7	T7-P7	P7-O1	FP1-F3	F3-C3	C3-P3	P3-O1	FZ-CZ	CZ-PZ
0	31.452991	-100.610501	-2.930403	4.102564	9.181929	-23.247863	-42.393162	-9.963370	5.665446	-57.631258
1	40.048840	-100.610501	-12.698413	12.698413	0.195360	-24.420024	-38.485958	3.711844	2.539683	-52.551893
2	36.141636	-87.716728	-18.949939	26.373626	-0.586081	-27.155067	-32.625153	18.559219	-2.148962	-37.704518
3	43.565324	-84.981685	-2.539683	16.214896	-3.711844	-25.982906	-22.075702	25.982906	-10.354090	-31.452991
4	57.631258	-96.312576	4.884005	20.122100	-6.446886	-29.890110	-15.433455	40.830281	-14.652015	-23.247863
...
3426811	-113.504274	82.246642	113.504274	-148.669109	-6.837607	-81.465201	44.737485	-22.466422	16.605617	-22.075702
3426812	-104.126984	32.625153	137.338217	-134.993895	-63.882784	-22.075702	35.360195	-17.777778	19.731380	-20.122100
3426813	-92.014652	-3.321123	113.504274	-97.875458	-39.267399	-37.704518	20.122100	-22.466422	28.717949	-15.824176
3426814	-68.571429	-24.029304	86.153846	-74.432234	-39.267399	-19.731380	16.214896	-39.658120	28.327228	-17.387057
3426815	-34.578755	-56.068376	96.312576	-74.041514	-74.432234	33.406593	23.638584	-51.770452	24.420024	-18.559219

3426816 rows x 22 columns

3. Data Splitting:

```
In [25]: # Define the features (X) and labels (y)
X = final_df.drop('seizure', axis=1)
y = final_df['seizure']
```

```
In [26]: # Split the data into training, validation, and test sets
X_train, X_temp, y_train, y_temp = train_test_split(X, y, test_size=0.3, random_state=42)
X_val, X_test, y_val, y_test = train_test_split(X_temp, y_temp, test_size=0.5, random_state=42)
```

```
In [27]: # Print the shapes of the resulting sets
print("Training set shape:", X_train.shape)
print("Validation set shape:", X_val.shape)
print("Test set shape:", X_test.shape)
```

```
Training set shape: (30, 2)
Validation set shape: (7, 2)
Test set shape: (7, 2)
```

```
In [28]: # Combine seizure and non-seizure data
df_seizure['seizure'] = 1
df_nonseizure['seizure'] = 0
```

```
In [29]: final_df = pd.concat([df_seizure, df_nonseizure], ignore_index=True)
```

```
In [30]: # Define the features (X) and labels (y)
X = final_df.drop('seizure', axis=1).values
y = final_df['seizure'].values
```

```
In [35]: # Split the data into training, validation, and test sets
X_train, X_temp, y_train, y_temp = train_test_split(X, y, test_size=0.3, random_state=42)
X_val, X_test, y_val, y_test = train_test_split(X_temp, y_temp, test_size=0.5, random_state=42)
```

```
In [36]: # Reshape the data for LSTM input (assuming time series data)
X_train = X_train.reshape((X_train.shape[0], 1, X_train.shape[1]))
X_val = X_val.reshape((X_val.shape[0], 1, X_val.shape[1]))
X_test = X_test.reshape((X_test.shape[0], 1, X_test.shape[1]))
```

```
In [39]: #CNN
```

```
In [59]: # Split data into features (X) and labels (y)
X_cnn = final_df.drop('seizure', axis=1)
y_cnn = final_df['seizure']

# Split the data into training, validation, and test sets
X_train_cnn, X_temp_cnn, y_train_cnn, y_temp_cnn = train_test_split(X_cnn, y_cnn, test_size=0.3, random_state=42)
X_val_cnn, X_test_cnn, y_val_cnn, y_test_cnn = train_test_split(X_temp_cnn, y_temp_cnn, test_size=0.5, random_state=42)

# Reshape data for CNN [samples, timesteps, features]
X_train_cnn = X_train_cnn.values.reshape((X_train_cnn.shape[0], X_train_cnn.shape[1], 1))
X_val_cnn = X_val_cnn.values.reshape((X_val_cnn.shape[0], X_val_cnn.shape[1], 1))
X_test_cnn = X_test_cnn.values.reshape((X_test_cnn.shape[0], X_test_cnn.shape[1], 1))
```

4. Model Selection:

Using BI-LTSM

```
In [41]: # Create a Bi-LSTM model
model_dropout = Sequential()
model_dropout.add(Bidirectional(LSTM(64, activation='relu'), input_shape=(1, X_train.shape[2])))
model_dropout.add(Dropout(0.5))
model_dropout.add(Dense(1, activation='sigmoid'))
```

Using CNN

```
In [60]: model_cnn = Sequential([
    Conv1D(filters=64, kernel_size=3, activation='relu', input_shape=(X_train_cnn.shape[1], 1)),
    MaxPooling1D(pool_size=2),
    Flatten(),
    Dense(50, activation='relu'),
    Dropout(0.3),
    Dense(1, activation='sigmoid')
])
```

5. Model Training:

BI-LTSM

```
In [43]: # Compile the model
model_dropout.compile(optimizer='adam', loss='binary_crossentropy', metrics=['accuracy'])

# Train the model
history_dropout = model_dropout.fit(X_train, y_train, epochs=10, batch_size=32, validation_data=(X_val, y_val))
```

```

Epoch 1/10
74962/74962 [=====] - 101s 1ms/step - loss: 0.3788 - accuracy: 0.8549 - val_loss: 0.349
5 - val_accuracy: 0.8597
Epoch 2/10
74962/74962 [=====] - 99s 1ms/step - loss: 0.3616 - accuracy: 0.8575 - val_loss: 0.3470
- val_accuracy: 0.8596
Epoch 3/10
74962/74962 [=====] - 98s 1ms/step - loss: 0.3600 - accuracy: 0.8578 - val_loss: 0.3448
- val_accuracy: 0.8597
Epoch 4/10
74962/74962 [=====] - 98s 1ms/step - loss: 0.3589 - accuracy: 0.8580 - val_loss: 0.3439
- val_accuracy: 0.8603
Epoch 5/10
74962/74962 [=====] - 98s 1ms/step - loss: 0.3588 - accuracy: 0.8582 - val_loss: 0.3450
- val_accuracy: 0.8597
Epoch 6/10
74962/74962 [=====] - 98s 1ms/step - loss: 0.3586 - accuracy: 0.8585 - val_loss: 0.3431
- val_accuracy: 0.8594
Epoch 7/10
74962/74962 [=====] - 98s 1ms/step - loss: 0.3583 - accuracy: 0.8585 - val_loss: 0.3430
- val_accuracy: 0.8604
Epoch 8/10
74962/74962 [=====] - 102s 1ms/step - loss: 0.3583 - accuracy: 0.8587 - val_loss: 0.345
8 - val_accuracy: 0.8599
Epoch 9/10
74962/74962 [=====] - 98s 1ms/step - loss: 0.3583 - accuracy: 0.8588 - val_loss: 0.3450
- val_accuracy: 0.8611
Epoch 10/10
74962/74962 [=====] - 99s 1ms/step - loss: 0.3585 - accuracy: 0.8588 - val_loss: 0.3451
- val_accuracy: 0.8609

```

```

In [44]: # Predictions on the validation set
y_val_pred = model_dropout.predict(X_val)
y_val_pred_binary = np.round(y_val_pred)

```

```

16064/16064 [=====] - 8s 489us/step

```

CNN

```
In [61]: # Compile the model
model_cnn.compile(optimizer='adam', loss='binary_crossentropy', metrics=['accuracy'])

# Train the model
history_cnn = model_cnn.fit(X_train_cnn, y_train_cnn, epochs=10, batch_size=64, validation_data=(X_val_cnn, y_val_cnn))

Epoch 1/10
37481/37481 [=====] - 52s 1ms/step - loss: 0.3627 - accuracy: 0.8555 - val_loss: 0.3408
- val_accuracy: 0.8614
Epoch 2/10
37481/37481 [=====] - 50s 1ms/step - loss: 0.3488 - accuracy: 0.8592 - val_loss: 0.3368
- val_accuracy: 0.8625
Epoch 3/10
37481/37481 [=====] - 50s 1ms/step - loss: 0.3442 - accuracy: 0.8603 - val_loss: 0.3313
- val_accuracy: 0.8624
Epoch 4/10
37481/37481 [=====] - 50s 1ms/step - loss: 0.3415 - accuracy: 0.8612 - val_loss: 0.3292
- val_accuracy: 0.8628
Epoch 5/10
37481/37481 [=====] - 50s 1ms/step - loss: 0.3402 - accuracy: 0.8614 - val_loss: 0.3287
- val_accuracy: 0.8632
Epoch 6/10
37481/37481 [=====] - 49s 1ms/step - loss: 0.3395 - accuracy: 0.8615 - val_loss: 0.3296
- val_accuracy: 0.8626
Epoch 7/10
37481/37481 [=====] - 49s 1ms/step - loss: 0.3385 - accuracy: 0.8617 - val_loss: 0.3257
- val_accuracy: 0.8632
Epoch 8/10
37481/37481 [=====] - 49s 1ms/step - loss: 0.3381 - accuracy: 0.8618 - val_loss: 0.3249
- val_accuracy: 0.8639
Epoch 9/10
37481/37481 [=====] - 49s 1ms/step - loss: 0.3377 - accuracy: 0.8619 - val_loss: 0.3260
- val_accuracy: 0.8631
Epoch 10/10
37481/37481 [=====] - 50s 1ms/step - loss: 0.3374 - accuracy: 0.8618 - val_loss: 0.3267
- val_accuracy: 0.8631
```

6. Model Evaluation:

BI-LTSM

```
In [45]: # Evaluate performance on the validation set
accuracy_val = accuracy_score(y_val, y_val_pred_binary)
confusion_matrix_val = confusion_matrix(y_val, y_val_pred_binary)
classification_report_val = classification_report(y_val, y_val_pred_binary)

print("Validation Set Performance:")
print(f"Accuracy: {accuracy_val:.4f}")
print("Confusion Matrix:")
print(confusion_matrix_val)
print("Classification Report:")
print(classification_report_val)
```

Validation Set Performance:

Accuracy: 0.8609

Confusion Matrix:

```
[[ 3767  71380]
 [   120 438755]]
```

Classification Report:

	precision	recall	f1-score	support
0	0.97	0.05	0.10	75147
1	0.86	1.00	0.92	438875
accuracy			0.86	514022
macro avg	0.91	0.52	0.51	514022
weighted avg	0.88	0.86	0.80	514022

CNN

```
In [62]: # Evaluate the model on the test set
test_loss_cnn, test_accuracy_cnn = model_cnn.evaluate(X_test_cnn, y_test_cnn)
print(f"Test accuracy: {test_accuracy_cnn * 100:.2f}%")
```